

In [1]:

```

1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn import preprocessing, svm
6 from sklearn.model_selection import train_test_split
7 from sklearn.linear_model import LinearRegression

```

In [2]:

```

1 df=pd.read_csv(r"C:\Users\HP\OneDrive\Documents\fiat500_VehicleSelection_Dataset.csv")
2 df

```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611111
1	2	pop	51	1186	32500	1	45.666359	12.244444
2	3	sport	74	4658	142228	1	45.503300	11.416667
3	4	lounge	51	2739	160000	1	40.633171	17.633333
4	5	pop	73	3074	106880	1	41.903221	12.494444
...
1533	1534	sport	51	3712	115280	1	45.069679	7.700000
1534	1535	lounge	74	3835	112000	1	45.845692	8.666667
1535	1536	pop	51	2223	60457	1	45.481541	9.416667
1536	1537	lounge	51	2557	80750	1	45.000702	7.683333
1537	1538	pop	51	1766	54276	1	40.323410	17.566667

1538 rows × 9 columns

In [3]:

```

1 df=df[['engine_power','age_in_days']]
2 df.columns=['ep','aid']

```

In [4]:

```
1 df.describe()
```

Out[4]:

	ep	aid
count	1538.000000	1538.000000
mean	51.904421	1650.980494
std	3.988023	1289.522278
min	51.000000	366.000000
25%	51.000000	670.000000
50%	51.000000	1035.000000
75%	51.000000	2616.000000
max	77.000000	4658.000000

In [5]:

```
1 df.head(10)
```

Out[5]:

	ep	aid
0	51	882
1	51	1186
2	74	4658
3	51	2739
4	73	3074
5	74	3623
6	51	731
7	51	1521
8	73	4049
9	51	3653

In [6]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0    ep      1538 non-null   int64  
 1    aid      1538 non-null   int64  
dtypes: int64(2)
memory usage: 24.2 KB
```

In [7]:

```
1 df.fillna(method='ffill',inplace=True)
```

C:\Users\HP\AppData\Local\Temp\ipykernel_22436\4116506308.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.fillna(method='ffill',inplace=True)
```

In [8]:

```
1 x=np.array(df['ep']).reshape(-1,1)
2 y=np.array(df['aid']).reshape(-1,1)
```

In [9]:

```
1 df.dropna(inplace=True)
```

C:\Users\HP\AppData\Local\Temp\ipykernel_22436\1379821321.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)


```
df.dropna(inplace=True)
```

In [10]:

```
1 #Exploring the data scatter_plotting the data scatter
2 sns.lmplot(x = "ep", y = "aid", data = df, order = 2, ci = None)
```

Out[10]:

<seaborn.axisgrid.FacetGrid at 0x20e4319bb50>




In [11]:

```
1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
2 regr=LinearRegression()
3 regr.fit(x_train,y_train)
4 print(regr.score(x_test,y_test))
```

0.08266478450766102

In [12]:

```
1 y_pred=regr.predict(x_test)
2 plt.scatter(x_test,y_test,color='b')
3 plt.plot(x_test,y_pred,color='k')
4 plt.show()
```




In [13]:

```
1 df100=df[:][:100]
2 sns.lmplot(x='ep',y='aid',data=df100,order=1,ci=None)
3
```

Out[13]:

<seaborn.axisgrid.FacetGrid at 0x20e28fa1ab0>



In [14]:


```

1 df100.fillna(method='ffill',inplace=True)
2 X=np.array(df100['ep']).reshape(-1,1)
3 y=np.array(df100['aid']).reshape(-1,1)
4 df100.dropna(inplace=True)
5 X_train,x_test,y_train,y_test=train_test_split(X,y,test_size=0.25)
6 regr=LinearRegression()
7 regr.fit(X_train,y_train)
8 print(regr.score(x_test,y_test))
9 print("Regression: ",regr.score(x_test,y_test))
10 y_pred=regr.predict(x_test)
11 plt.scatter(x_test,y_test,color='g')
12 plt.plot(x_test,y_pred,color='k')
13 plt.show()

```

0.08993389468303326

Regression: 0.08993389468303326



In [15]:

```

1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 model=LinearRegression()
4 model.fit(X_train,y_train)
5 y_pred=model.predict(x_test)
6 r2=r2_score(y_test,y_pred)
7 print("R2_score: ",r2)

```

R2_score: 0.08993389468303326

Conclusion:

Dataset we have taken is poor for linear model but with the smaller data works well with linear model

In [16]:

```
1 from sklearn.linear_model import LinearRegression
2 from sklearn.model_selection import train_test_split
3 feature=df.columns[0:3]
4 target=df.columns[-1]
5 x=df[feature].values
6 y=df[target].values
7 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25)
8 lr = LinearRegression()
9 lr.fit(x_train,y_train)
10 print(lr.score(x_test,y_test))
11 print(lr.score(x_train,y_train))
```

1.0

1.0

In [17]:

```
1 from sklearn.linear_model import Ridge,RidgeCV,Lasso,LassoCV
2 ridge = Ridge(alpha=10)
3 ridge.fit(x_train,y_train)
4 train_score_ridge=ridge.score(x_train,y_train)
5 test_score_ridge=ridge.score(x_test,y_test)
6 print('\n Ridge method\n')
7 print('The score of ridge method is {}'.format(train_score_ridge))
8 print('The score of ridge method is {}'.format(test_score_ridge))
```

Ridge method

The score of ridge method is 1.0

The score of ridge method is 1.0

In [18]:

```
1 lasso = Lasso(alpha=10)
2 lasso.fit(x_train,y_train)
3 train_score_lasso=lasso.score(x_train,y_train)
4 test_score_lasso=lasso.score(x_test,y_test)
5 print('\n Lasso method\n')
6 print('The score of lasso method is {}'.format(train_score_lasso))
7 print('The score of lasso method is {}'.format(test_score_lasso))
```

Lasso method

The score of lasso method is 0.9999999999632855

The score of lasso method is 0.9999999999632833

In [19]:

```
1 lasso_cv= LassoCV(alphas=[0.2,0.03,0.004,0.0001,1,20]).fit(x_train,y_train)
2 train_score_lasso_cv=lasso_cv.score(x_train,y_train)
3 test_score_lasso_cv=lasso_cv.score(x_test,y_test)
4 print('\n LassoCV method\n')
5 print('The score of Lasso method is {}'.format(train_score_lasso_cv))
6 print('The score of Lasso method is {}'.format(test_score_lasso_cv))
```

LassoCV method

The score of Lasso method is 1.0

The score of Lasso method is 1.0

In [20]:

```
1 ridge_cv=RidgeCV(alphas=[1,2.3,0.2,0.3,0.4,0.5,0.6]).fit(x_train,y_train)
2 print("\n RidgeCV Method\n")
3 print("The score of Ridge method is {}".format(ridge_cv.score(x_train,y_train)))
4 print("The score of Ridge method is {}".format(ridge_cv.score(x_test,y_test)))
```


RidgeCV Method

The score of Ridge method is 0.9999999996466192

The score of Ridge method is 0.9999999996465979


In [21]:

```
1 plt.figure(figsize=(10,10))
2 plt.plot(feature,ridge.coef_,alpha=0.5,marker='*',markersize=5,linestyle='None',color='blue')
3 plt.plot(feature,lr.coef_,alpha=0.7,marker='o',markersize=7,color='green',linestyle='None')
4 plt.xticks(rotation=90)
5 plt.title("Comparision between Ridge and linear regression")
6 plt.legend()
7 plt.show()
```



In [22]:

```
1 plt.figure(figsize=(10,10))
2 plt.plot(feature,ridge.coef_,alpha=0.5,marker='*',markersize=5,linestyle='None',col
3 plt.plot(feature,ridge.coef_,alpha=0.6,marker='d',markersize=6,linestyle='None',col
4 plt.plot(feature,lr.coef_,alpha=0.7,marker='o',markersize=7,color='green',linestyle
5 plt.xticks(rotation=90)
6 plt.title("Comparision between Ridge and linear regression")
7 plt.legend()
8 plt.show()
```



In []:

1